

Topics in Computational Inference

David M. Allen
University of Kentucky

November 4, 2018

12 Calling C and C++ from R

The definitive reference for calling C or C++ from R is Chapter 5 of [Writing R Extensions](#). The .C method is presented here. This method is also favored by Eubank and Kupresanin [1] and Matloff [2].

Rcpp is a more elaborate method, and is discussed in [Rcpp: Seamless R and C++ Integration](#).

Section 12.1 — Introduction

R is a wonderful programming language. It can leap long matrix expressions in a single bound. However *R* can be slow for situations where frequent access to individual elements of a matrix and loops are required. The *C* and *C++* languages are often more efficient in these situations. A good strategy is to program such functions in *C* or *C++* and call them from *R*. This note gives a recipe for doing that.

A Distinction Between C and C++

The function names must be unique in a C program. This is a benefit when interfacing with R. C++ allows function overloading, *ie.* multiple functions with the same name. This requires *name mangling* in order to have unique identifiers. Also, class methods don't behave exactly like ordinary functions. Therefore, I give a C example and a C++ example.

Two Situations

I consider two situations:

- Code is being developed from the beginning with the objective of interfacing it with *R*. Consider using *C*, as interfacing it is simpler.
- There is an existing C++ class that is well tested and you want to use it without modification.

The Fibonacci Series

The Fibonacci series f_0, f_1, \dots is a series where $f_0 = 0, f_1 = 1$, and all subsequent terms are the sum of the two preceding terms. Calculation of the first n terms involves a loop and is easily implemented in C.

The C Code

The demonstration *C* function takes the desired number of terms as input and returns the series. The file that contains the function must have an extension `.c`. The function must return `void`, and the arguments must be pointers. The code for the Fibonacci series is in Listing 12.1.

Listing 12.1: Fibonacci.c

```
#include <R.h>
// typically n would not be a pointer
void Fibonacci(int* n, int* f)
{
    if(*n < 2)
        *n = -1;
    else
    {
        f[0] = 0;
        f[1] = 1;
        for(int i = 2; i < *n; i++)
            f[i] = f[i-1] + f[i-2];
    }
    return;
}
```


Compilation

The code is in `Fibonacci.c` and is compiled from the command line by

```
R CMD SHLIB Fibonacci.c
```

On Linux or Unix, this will produce a file `Fibonacci.so`. The “so” stands for *shared object*. On Windows this will produce a file, `Fibonacci.dll`. The “dll” stands for *dynamic linked library*.

Calling the Function from *R*

The calling *R* program must have

```
dyn.load("Fibonacci.so")
```

before the function is called. The function is called by `.C(args)` where `args` is the name of the function enclosed in double quotes, followed by the function arguments cast to their types. The code is in Listing [12.2](#).

Listing 12.2: Fibonacci.R

```
dyn.load("Fibonacci.so")

n      <- 10
# create memory for the result
f      <- vector(mode = "integer", length = n);
# f     <- rep(0, n)   # this would also work
result <- .C("Fibonacci", as.integer(n), as.integer(f))
# result contains a list with an element for
# each argument
print( result[[2]] )
```

The Convolve Function

This example takes a C++ class and *wraps it*, so that it looks like a group of C functions. The vehicle for illustration is the convolve function.

Convolve is the verb form of convolution. We might say we convolve vectors a and b to produce a convolution. In statistics the concept comes into play when finding the distribution of a sum of random variables. With reference to the next slide, the convolution of vectors a and b is a vector containing the sums within the blue bars.


$$a_1 \times b_1$$

$$a_1 \times b_2$$

$$a_1 \times b_3$$

$$a_1 \times b_4$$

$$a_2 \times b_1$$

$$a_2 \times b_2$$

$$a_2 \times b_3$$

$$a_2 \times b_4$$

$$a_3 \times b_1$$

$$a_3 \times b_2$$

$$a_3 \times b_3$$

$$a_3 \times b_4$$

The Source Code

The class is such that constructor sets the size of the input vectors. Any number of vectors of that size can be processed with a single object of the class.

The header code is in Listing 12.3. The C++ code is in Listing 12.4.

Listing 12.3: Convolve.h

```
#ifndef CONVOLVE
#define CONVOLVE

// design is such that evaluate could be called
// multiple times for same na & nb
class Convolve
{
    int      na;
    int      nb;
public:
    Convolve(int na, int nb);
    ~Convolve();
    void      evaluate(const double* a, const double* b,
                      double* ab);
};
#endif
```

Listing 12.4: Convolve.cpp

```
#include <cstring> // memset is here
#include "Convolve.h"

Convolve::Convolve(int na, int nb)
{
    this->na = na; this->nb = nb;
}

Convolve::~Convolve(){}

void Convolve::evaluate(const double* a, const double* b,
                        double* ab)
{
    memset(ab, 0, (na + nb - 1)*sizeof(double));
    for(int i = 0; i < na; i++)
        for(int j = 0; j < nb; j++)
            ab[i + j] += a[i] * b[j];
}
```


Unit Testing

This is a brief diversion. Often when one builds a class, it is with the idea that it will be a module in a large system or program. However there should be a program that is designed to do nothing but test all the features of the class. This process of testing classes in isolation is called *unit testing*.

Code that exercises the Convolve class is in Listing [12.5](#).

Listing 12.5: driver.cpp

```
#include <cstdio>
#include "Convolve.cpp"

int main(int argc, char *argv[])
{
    const int n = 6;
    double a[n];
    double b[n];

    for(int i = 0; i<n; i++)
    {
        a[i] = 1.0/n;
        b[i] = 1.0/n;
    }

    Convolve* con;
    double* ab = new double[n+n-1];
    con = new Convolve(n, n);
    con->evaluate(a, b, ab);
}
```

```
fprintf(stdout, "Distribution of sum\n");  
fprintf(stdout, "%10s%10s\n", "value", "prob");  
for(int i = 0; i<(n+n-1); i++)  
    fprintf(stdout, "%10i%10.4lf\n", i+2, ab[i]);  
return 0;  
}
```

Compiling the Driver

The following are instructions compiling and running the program using the *GNU Compiler Collection (GCC)*. If you use another compiler, see its instructions for compiling and running programs from source code. Open a terminal window and change to the directory containing the files. Issue the command

```
g++ driver.cpp -o driver
```

On Windows use `-o driver.exe`.

Run the Driver

Run the program with the command

```
./driver
```

Windows users could omit the “./”. Check the output with what it should be. Correct any errors and repeat.

The R Interface File

For reasons mentioned earlier, C++ can't be directly loaded into R. The solution is to wrap C++ functions with C functions that follow conventions required by R. The interface or *wrapper* code for the current example is in Listing [12.6](#).

Listing 12.6: R-Convolve.cpp

```
#include <cstdio>
#include <cstdlib>
#include <R.h>
#include "Convolve.h"

static Convolve* conv;
extern "C"
{

void setup(int* na, int* nb) // arguments are pointers
{
    conv = new Convolve(*na, *nb);
}

void evaluate(double* a, double* b, double* ab)
{
    conv->evaluate(a, b, ab);
}
```

```
void closeout()
{
    delete conv; // calls destructor
}

} // end extern "C"
```


Compiling

Convolve.cpp and R-Convolve.cpp are compiled together from the command line by

```
R CMD SHLIB Convolve.cpp R-Convolve.cpp
```

This produces Convolve.so (or dll). The name is taken from the first file in the list.

Using the Shared Object File

The code that loads and uses Convolve.so is

Listing 12.7: Convolve.R

```
dyn.load("Convolve.so")

lengths <- function(na, nb)
{
  .C("setup", as.integer(length(a)),
                                     as.integer(length(b)))
  sprintf("For vectors of length %d and %d.\n", na, nb)
}

convolution <- function(a, b)
{
  ab <- rep(0.0, length(a)+length(b)-1)
  r <- .C("evaluate", as.numeric(a), as.numeric(b),
                                     as.numeric(ab))
  r[[3]]
}
```

```
}  
  
a <- rep(1,6)/6.0  
b <- a  
  
lengths(length(a), length(b))  
ab <- convolution(a, b)  
  
print(a)  
print(b)  
dist <- data.frame(values = 2:12, probs = ab)  
print(dist)
```

An Exercise

Exercise 12.1. In a C program you have variables

```
double a[] = {1, 2, 3, 6, 5, 4};  
int m = 2;  
int n = 3;
```

but pretend these are the result of long arduous computations. Write interface functions so that in the end you have a 2×3 matrix

$$\begin{bmatrix} 1 & 2 & 3 \\ 6 & 5 & 4 \end{bmatrix}$$

in R. Zip up all your files and email them to me.

References

- [1] Randall L. Eubank and Ana Kupresanin. *Statistical Computing in C++ and R*. The R Series. Chapman & Hall/CRC, 2012.
- [2] Norman Matloff. *The Art of R Programming*. San Francisco: No Starch Press, 2011.